

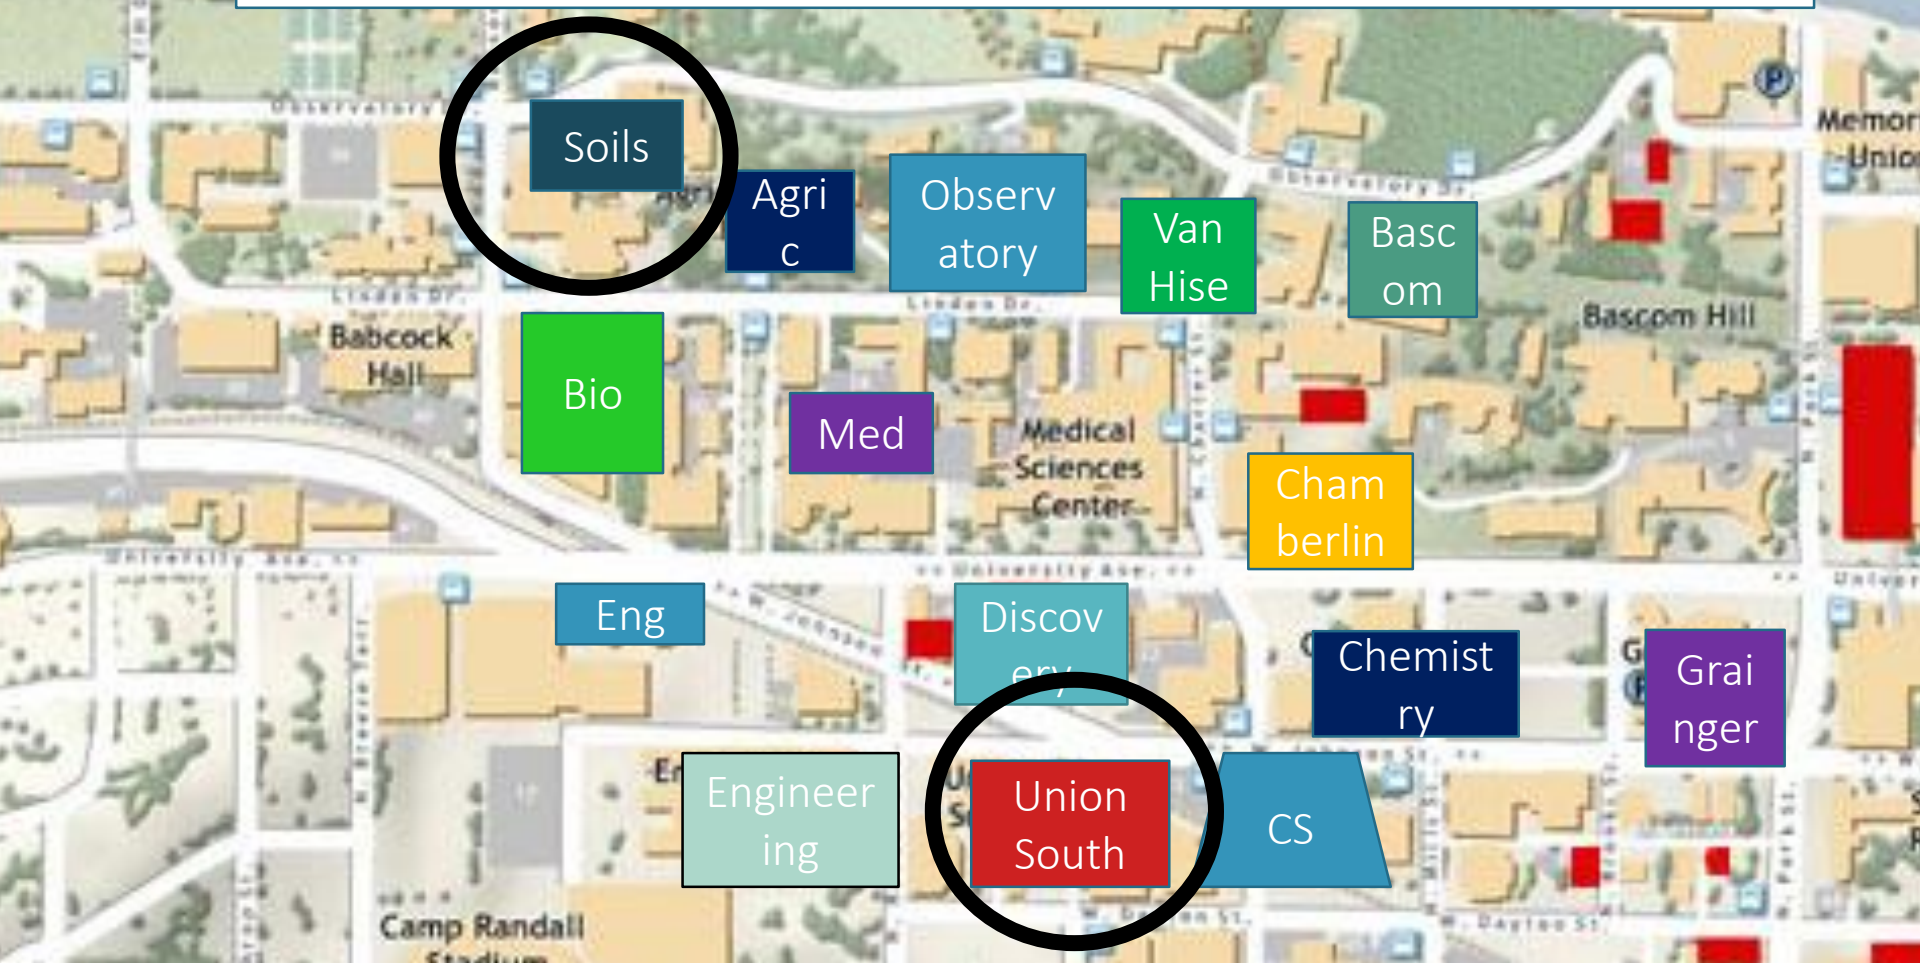
CS-540: Intro to Artificial Intelligence

SECTION 1

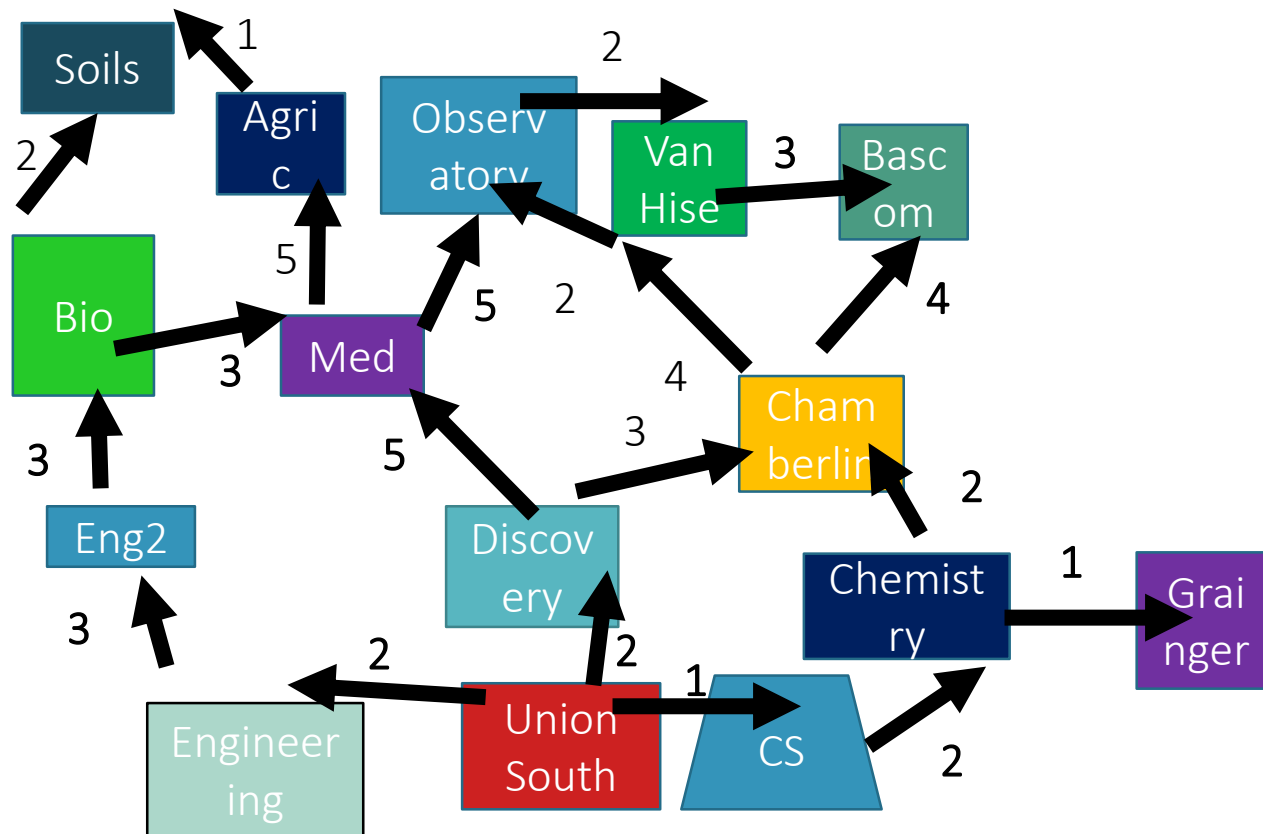
LECTURER: ERIN WINTER

Informed Search

Given the search techniques you currently know, how would you find the fastest route from Union South to Soil Sciences?



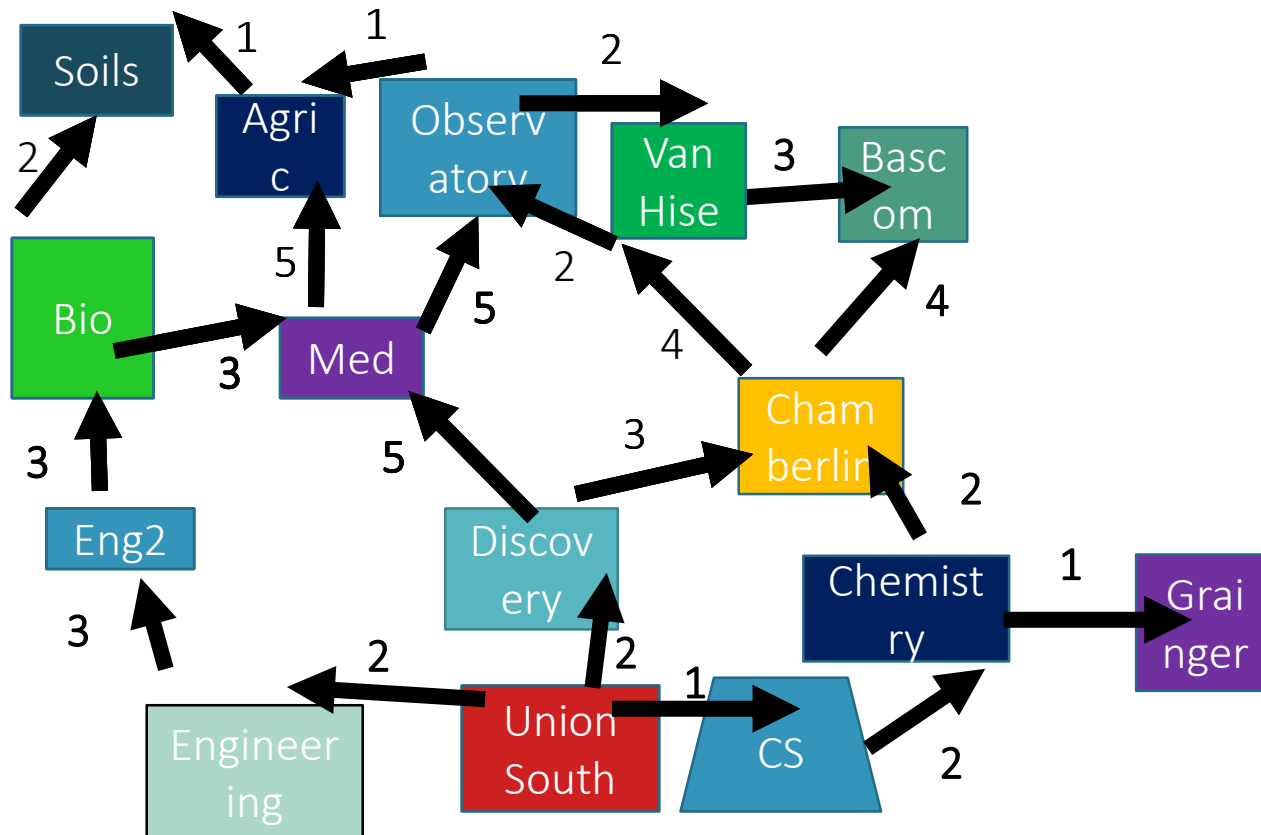
Because you have a **weighted, directed** graph Dijkstra's Search is your only suitable choice. It is also **optimal** and **complete**. Observe the first few states that Dijkstra's chooses to expand.



*Assume ties broken
alphabetically

*Allow state to appear multiple
times in priority queue

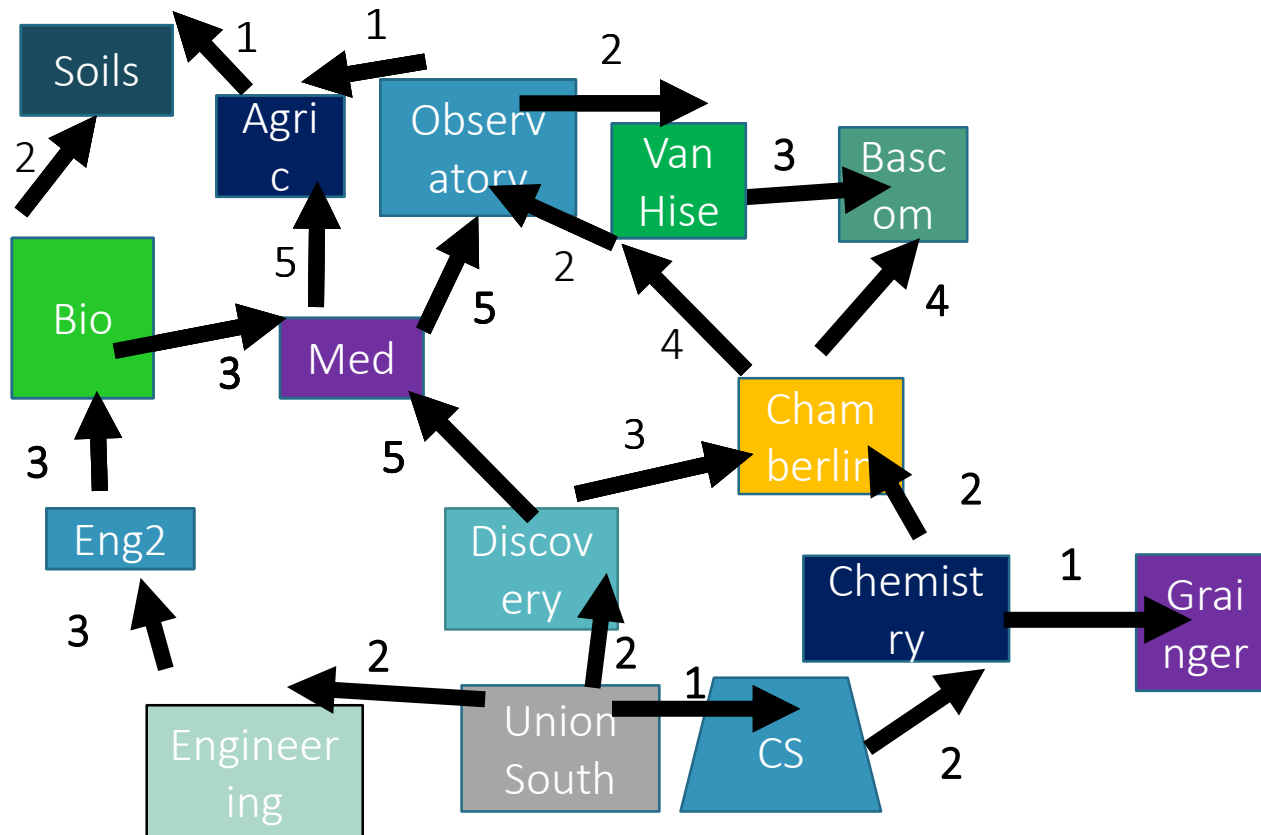
Frontier	Expanded
{US:0}	{}



*Assume ties broken
alphabetically

*Allow state to appear multiple
times in priority queue

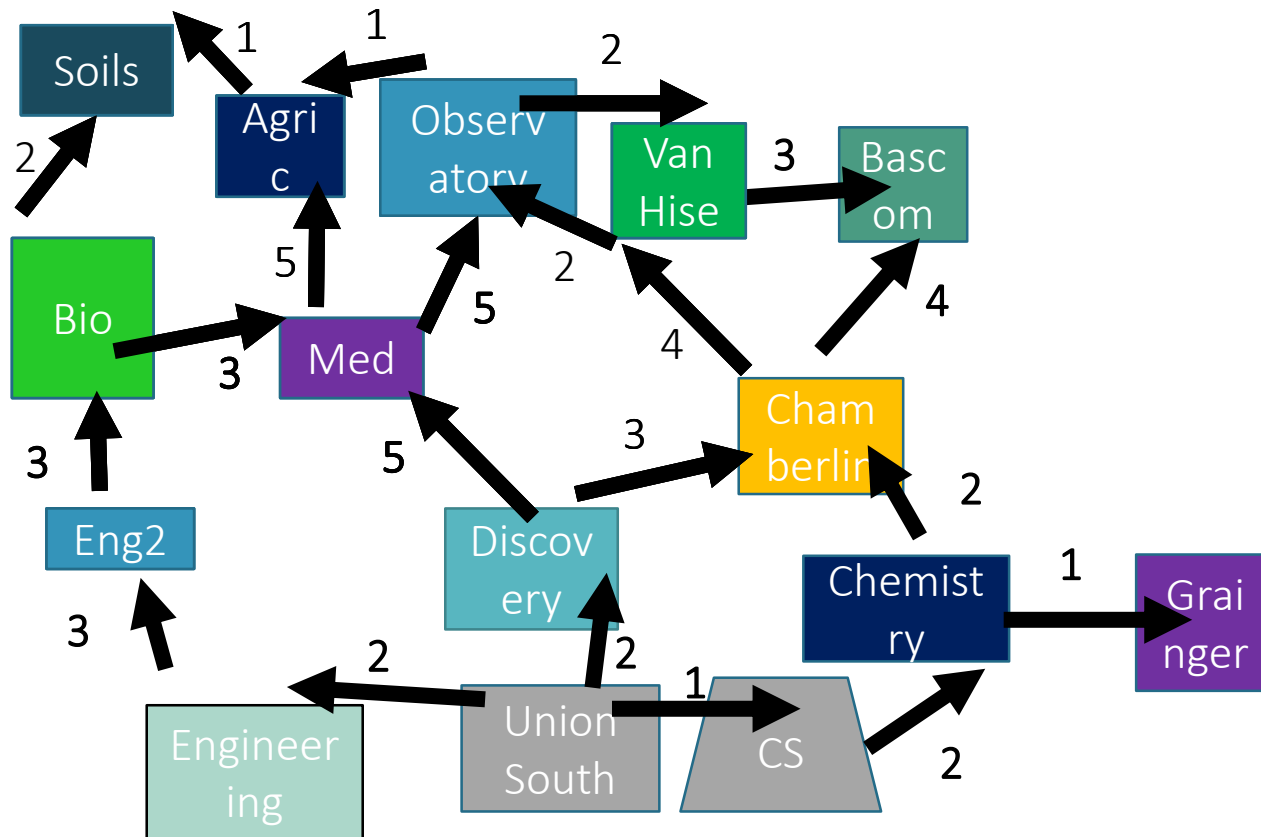
Frontier	Expanded
{US:0}	{}
{CS:1,DIS:2,ENG:2}	{US}



*Assume ties broken
alphabetically

*Allow state to appear multiple
times in priority queue

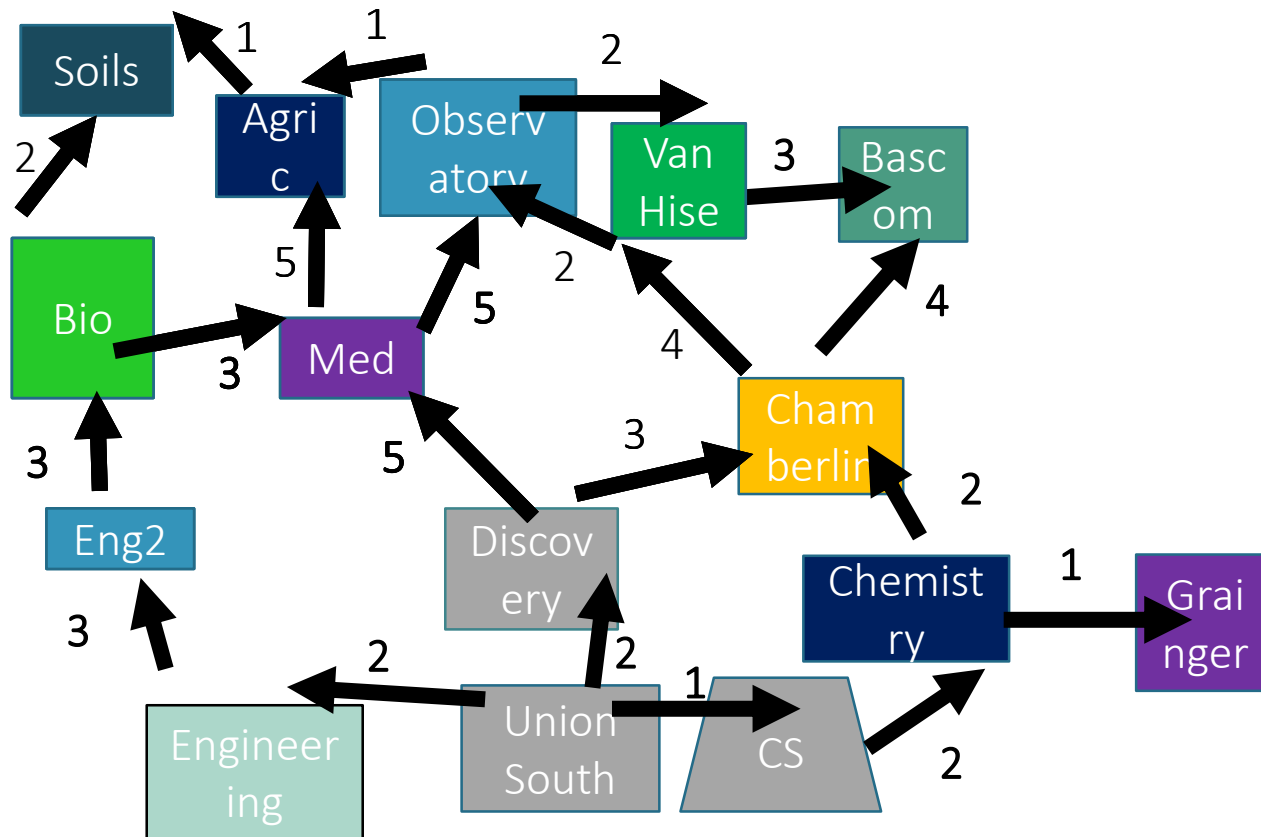
Frontier	Expanded
{US:0}	{}
{CS:1,DIS:2,ENG:2}	{US}
{DIS:2, ENG:2,CHEM:3}	{US,CS}



*Assume ties broken
alphabetically

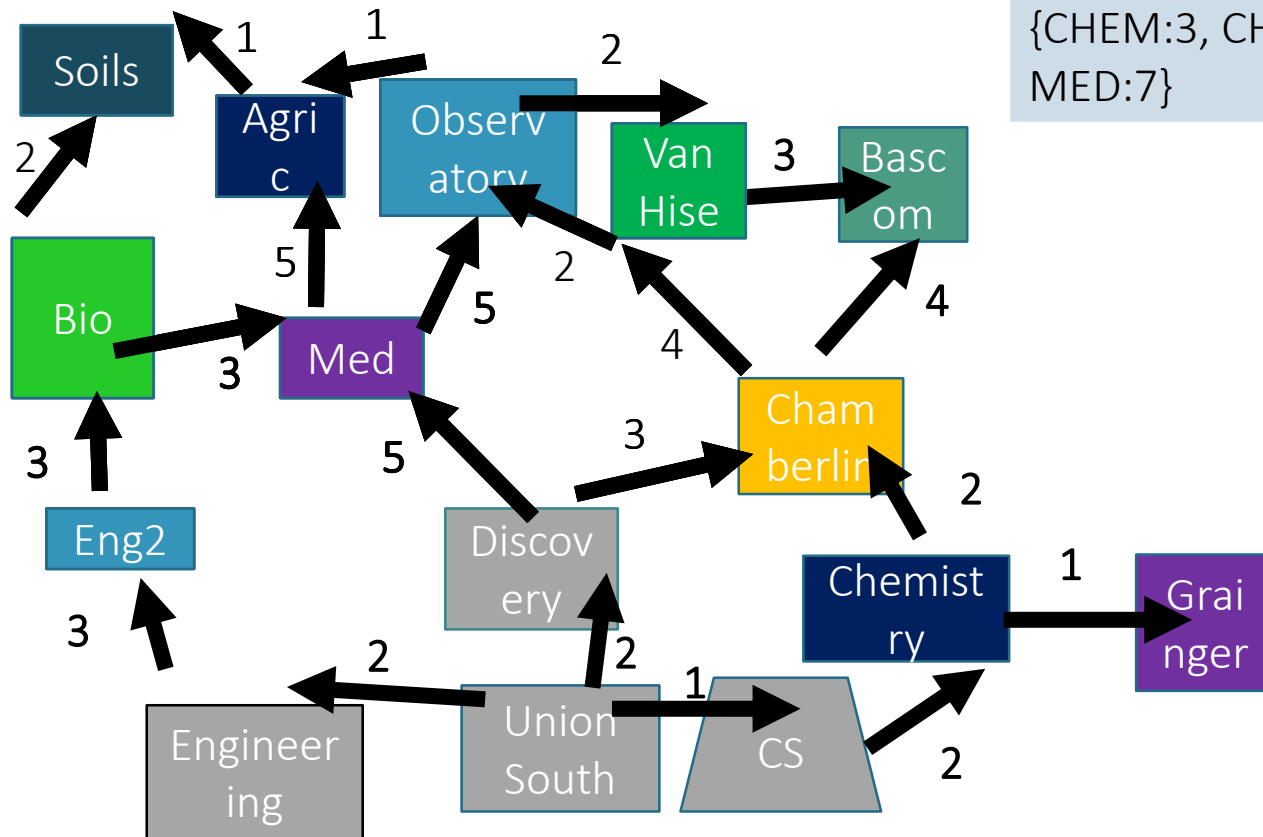
*Allow state to appear multiple
times in priority queue

Frontier	Expanded
{US:0}	{}
{CS:1,DIS:2,ENG:2}	{US}
{DIS:2, ENG:2,CHEM:3}	{US,CS}
{ENG:2,CHEM:3, CHA:5,MED:7}	{US,CS,DIS}



*Assume ties broken
alphabetically

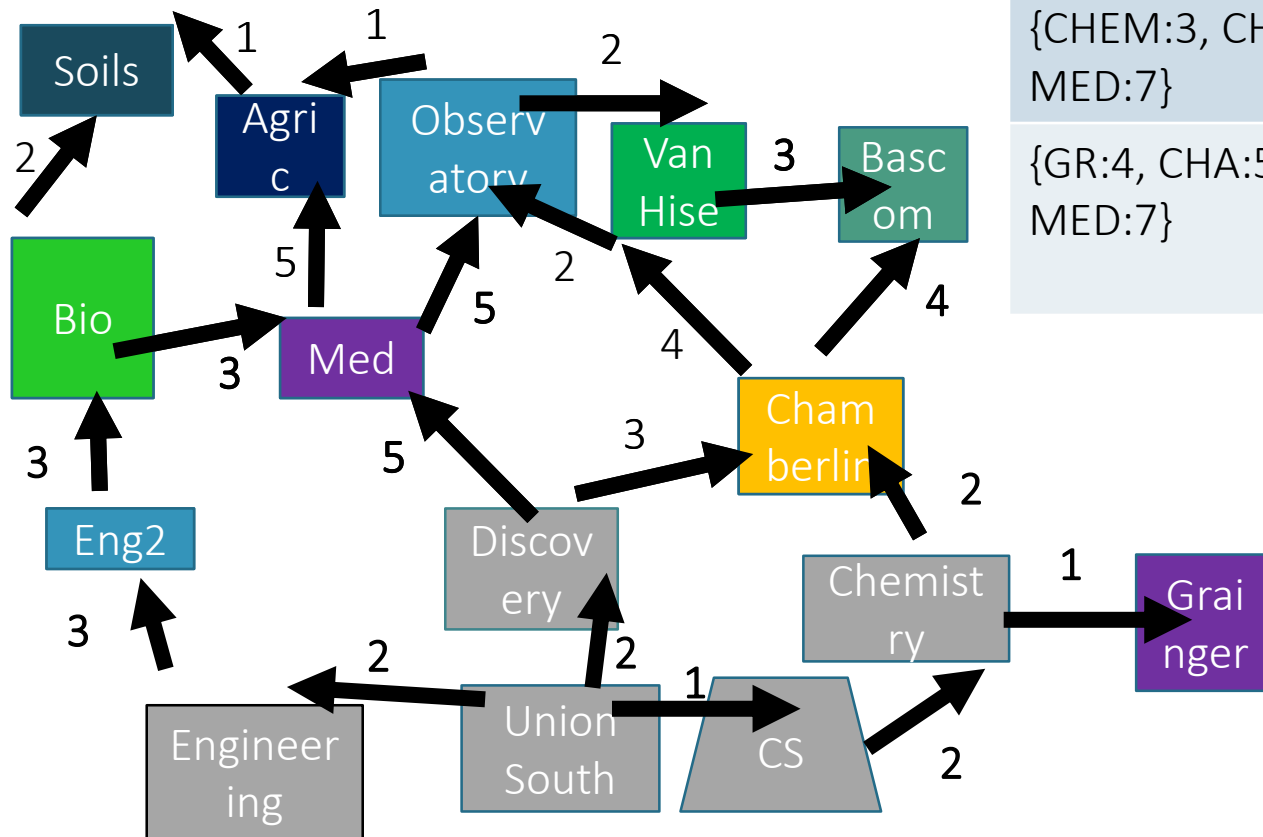
*Allow state to appear multiple
times in priority queue



Frontier	Expanded
{US:0}	{}
{CS:1,DIS:2,ENG:2}	{US}
{DIS:2, ENG:2,CHEM:3}	{US,CS}
{ENG:2,CHEM:3, CHA:5,MED:7}	{US,CS,DIS}
{CHEM:3, CHA:5, EN2:5, MED:7}	{US, CS, DIS, ENG}

*Assume ties broken alphabetically

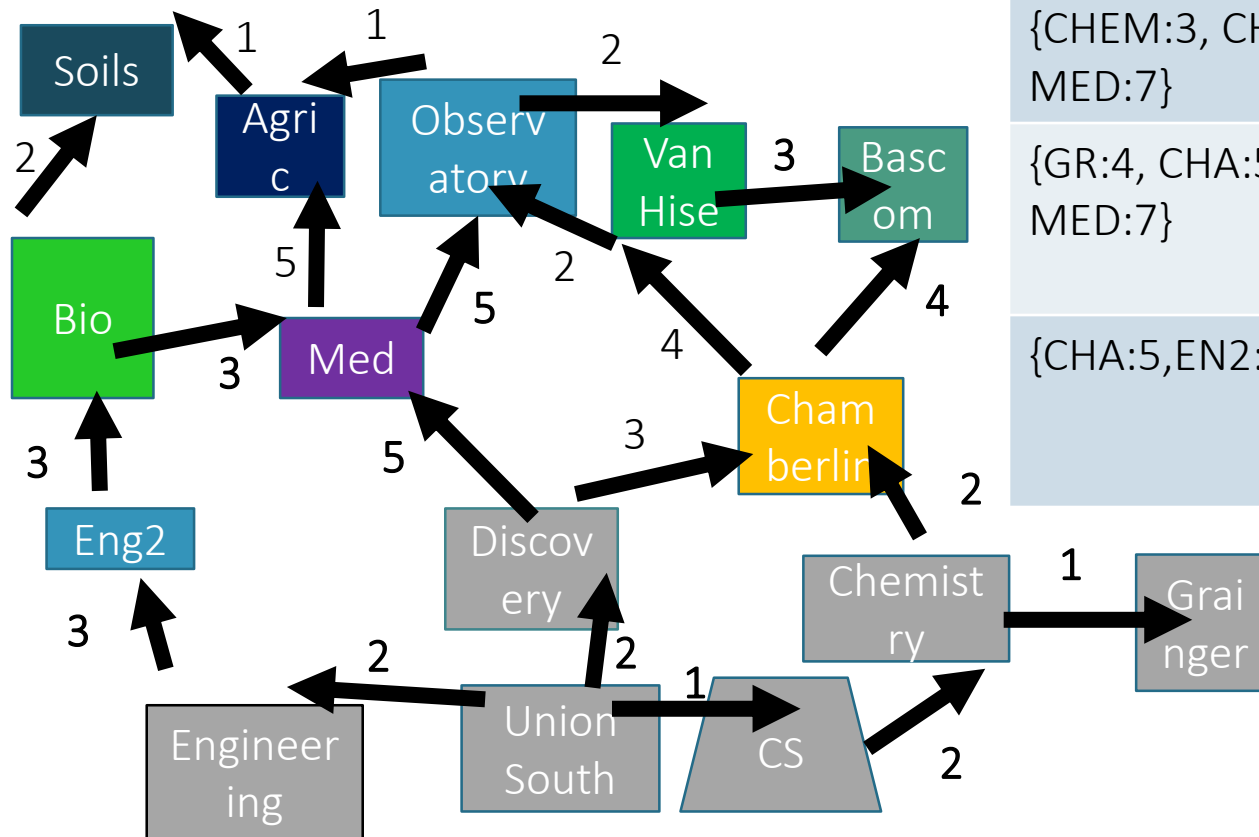
*Allow state to appear multiple times in priority queue



Frontier	Expanded
{US:0}	{}
{CS:1,DIS:2,ENG:2}	{US}
{DIS:2, ENG:2,CHEM:3}	{US,CS}
{ENG:2,CHEM:3, CHA:5,MED:7}	{US,CS,DIS}
{CHEM:3, CHA:5, EN2:5, MED:7}	{US, CS, DIS, ENG}
{GR:4, CHA:5, EN2:5, MED:7}	{US, CS,DIS, ENG, CHEM}

*Assume ties broken alphabetically

*Allow state to appear multiple times in priority queue



Frontier	Expanded
{US:0}	{}
{CS:1,DIS:2,ENG:2}	{US}
{DIS:2, ENG:2,CHEM:3}	{US,CS}
{ENG:2,CHEM:3,CHA:5,MED:7}	{US,CS,DIS}
{CHEM:3, CHA:5, EN2:5, MED:7}	{US, CS, DIS, ENG}
{GR:4, CHA:5, EN2:5, MED:7}	{US, CS,DIS, ENG, CHEM}
{CHA:5,EN2:5,MED:7}	{US,CS,DIS, ENG,CHEM ,GR}

What you know that Dijkstra's doesn't...

- The distance of each state relative to Soil Sciences!
- It doesn't make sense *in context* to investigate Grainger Hall or the Chemistry building as potential stops on the shortest route to Soil Sciences.
- How to reduce wasted effort? **A heuristic function that utilizes relative distance from the goal state.**

Big Idea of Informed Search

- Search problems are a pursuit of some goal state through a succession of intermediary states
- Instead of ordering states within the frontier by path cost, try to quantify the distance of a given state to a goal state using a **heuristic function**
- Heuristic functions are educated **guesses**, even with domain knowledge we can't expect them to be perfect

Heuristic Functions

Define a **heuristic** function, $h(n)$

- uses **domain-specific information** in some way
- is computable from the current state description
- it estimates
 - *the "goodness" of node n*
 - *how close node n is to a goal*
 - *the cost of minimal cost path from node n to a goal state*

Heuristic Function Properties

- $h(n) \geq 0$, for all nodes n *
- $h(n)$ close to 0 means we think n is close to a goal state
- $h(n)$ much greater than 0 means we think n is far from a goal state

*this property is especially important for ensuring correctness in algorithms that utilize heuristic functions

(Greedy) Best-First Search

- Implement a graph search with a **priority queue** frontier
- Prioritize nodes by lowest $h(n)$ value, that is, expand nodes estimated closest to the goal first
- Called a “greedy” algorithm because it makes shortsighted decisions based on what seems best at the given moment – selecting the lowest $h(n)$ value node from the frontier

Greedy Best-First Search


- Is it optimal?
- Is it complete?
- If it doesn't have both of these properties, it would be wiser to simply use Dijkstra's algorithm.

Greedy Best-First Search

$$f(n) = h(n)$$

of nodes tested: 0, expanded: 0

expnd. node	Frontier
	{S:8}




Greedy Best-First Search

$$f(n) = h(n)$$

of nodes tested: 1, expanded: 1

expnd. node	Frontier
	{S:8}
S not goal	{C:3,B:4,A:8}




Greedy Best-First Search

$$f(n) = h(n)$$

of nodes tested: 2, expanded: 2

expnd. node	Frontier
	{S:8}
S	{C:3,B:4,A:8}
C not goal	{G:0,B:4,A:8}



Greedy Best-First Search


$$f(n) = h(n)$$

of nodes tested: 3, expanded: 2

expnd. node	Frontier
	{S:8}
S	{C:3,B:4,A:8}
C	{G:0,B:4, A:8}
G	{B:4, A:8}

Optimal Path: S, B, G

Optimal Path Cost: 9



BFS Path: S,C,G

BFS Path Cost: 13

Learning from GBFS's Mistakes

- Greedy best-first search makes the unsound assumption that the heuristic function **always estimates distance to a goal state correctly**
- With an especially bad heuristic function, greedy best-first search can traverse the entire graph **$O(b^m)$** or worse, get on an infinite incorrect path like depth-first search
- Time and space complexity: **$O(b^m)$**

Attempt 2: A-Search

same computation as Dijkstra's




- Use as an evaluation function $f(n) = g(n) + h(n)$, where $g(n)$ is minimum cost path from start to current node n (as defined in UCS)
- Nodes in *the frontier* are ranked by the estimated cost of a solution, where $g(n)$ is the cost from the start node to node n , and $h(n)$ is the estimated cost from node n to a goal
- Think of it as a “first half” (from the start to node n using known edges) + “second half” (estimated cost from n to goal) computation

Evaluating A-Search

- **Smarter**, but still falls into the same of the same traps as Greedy Best-First Search with a bad heuristic
- Neither admissible nor complete.
- $O(b^m)$ time and space complexity.

Counterexample:

Because the heuristic function returns infinity on node E, E looks so bad that it will never be expanded



How to improve A-Search...

- Constrain the heuristic function!
- A-Search incorporates desirable information, cost to current node and the estimated cost from current node to the goal
- However, any heuristic function that **prevents expansion of nodes along the goal path** can never be optimal or complete.

Admissible Heuristics

- An admissible heuristic $h(n)$ never **overestimates** the cost from n to the goal
- So, an admissible heuristic applied to A-search $f(n) = g(n) + h(n)$ **can never overestimate the true solution path cost**

Admissible Heuristic Functions

Which of the following are admissible heuristics? Assume h^* is the true cost from n to the goal.

$$h(n) = h^*(n) \text{ YES}$$

$$h(n) = \max(2, h^*(n)) \text{ NO, what if } h^*(n) = 1$$

$$h(n) = \min(2, h^*(n)) \text{ YES, this never overestimates}$$

$$h(n) = h^*(n) - 2 \text{ NO, negative if } h^*(n) < 2$$

$$h(n) = \sqrt{h^*(n)} \text{ NO, square root of } h^*(n) \text{ is } > h^*(n) \text{ if } 0 < h^*(n) < 1$$

Comparing Heuristic Functions

- A heuristic function h_2 **dominates** h_1 if *for all* s

$$h_1(n) \leq h_2(n) \leq h^*(n)$$

- Want heuristic functions as close to h^* as possible, **but not over h^***
- Good heuristic functions can be slow, so sometimes it's better to use a weaker heuristic that expands more nodes in a given interval

A-Search + Admissible (or
Consistent) Heuristic =

A* Search

effective and used in practice

A* Algorithm

Frontier = {S} where S is the start node

Explored = {}

Loop do

if Frontier is empty **then return** failure

 pick node, n, with min $f(n) = h(n) + g(n)$ value from Frontier

if n is a goal state **then return** solution

for each each child, n', of n **do**

if n' is not in Explored or Frontier

then add n' to Frontier

else if n' in Frontier as m


if $g(m) \leq g(n')$ throw n' away

else add n' to Frontier and remove m

 Remove n from Frontier and add n to Explored

When Should A* Stop?

To ensure **optimality**, A* should terminate only when a goal state is *removed* from the priority queue



Same rule as for Dijkstra's

A* with $h(n) = 0$ *is Dijkstra's*

Consistent Heuristic Functions

A heuristic, h , is called **consistent** (aka **monotonic**) if, for every node n and every successor n' of n , the estimated cost of reaching the goal from n is *no greater than* the step cost of getting to n' plus the estimated cost of reaching the goal from n' :

$$c(n, n') \geq h(n) - h(n')$$

or, equivalently: $h(n) \leq c(n, n') + h(n')$

- Values of $f(n)$ along any path are *nondecreasing*
- **Consistency is a stronger condition than admissibility**

Admissible vs. Consistent Heuristics


A* using an **admissible** heuristic will give an optimal solution on an acyclic graphs (trees)

However, a **consistent** heuristic must be used with A* to guarantee optimality on graphs with cycles

Remember, all *consistent heuristics are admissible, but not all admissible heuristics are consistent*

Consistent Heuristics

Is this h consistent?



$h(C)=900$, $h(D)=1$, $c(C, D) = 1 < 900 - 1$, so h is *NOT* consistent

A* Search Practice

$$f(n) = g(n) + h(n)$$


of nodes tested: 0, expanded: 0

expnd. node	Frontier
	{S:0+8}

To test for admissibility, check that $h(n)$ never overestimates the true cost to the goal.

To check for consistency, check $h(n)$ is consistent AND $h(n) \leq c(n, n') + h(n')$ for all successor nodes n' .

Is h is admissible and/or consistent? Consistent




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 1, expanded: 1

expnd. node	Frontier
	{S:8}
S not goal	{A:1+8,B:5+4,C:8+3}




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 2, expanded: 2

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A not goal	{B:9,G:1+9+0,C:11, D:1+3+ ∞ ,E:1+7+ ∞ }




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 3, expanded: 3

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A	{B:9,G:10,C:11,D:∞,E:∞}
B not goal	{G:5+4+0, G:10 , C:11, D:∞,E:∞} replace




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 4, expanded: 3

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A	{B:9,G:10,C:11,D:∞,E:∞}
B	{G:9,C:11,D:∞,E:∞}
G goal	{C:11,D:∞,E:∞} not expanded




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 4, expanded: 3

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A	{B:9,G:10,C:11,D:∞,E:∞}
B	{G:9,C:11,D:∞,E:∞}
G	{C:11,D:∞,E:∞}



Fast, optimal, complete under same graph conditions as Dijkstra's

path: S,B,G
cost: 9

The Bad News About A^*

- A^* uses a lot of memory, ie. $O(\text{number of states})$
- For some search problems, A^* will fill the available system memory before finding a solution

How do we use less memory?

1. Sacrifice optimality – beam search
2. Sacrifice some time – iterative deepening A^*

Beam Search

Use an evaluation function $f(n) = h(n)$ as in Greedy Best-First search, *and* restrict the maximum size of the Frontier to a constant, k

- Only keep k best nodes as candidates for expansion, and throw away the rest
- More space efficient than Greedy Best-First Search, but may throw away a node on a solution path. Space complexity $O(km)$.
- Not complete
- Not optimal/admissible

Iterating Deepening A*

- Iterative-deepening A*
- Cutoff based on $f = g + h$ value rather than depth
- At each iteration do loop-avoiding DFS, not expanding any node with f -value that exceeds current threshold
- At each iteration increase the f -value threshold by setting it to the smallest f -value of any node that exceeded the cutoff in the previous iteration
- Complete
- Optimal / Admissible
- Linear space required

How to come up with admissible/consistent heuristics?

Heuristics are often defined by **relaxing the problem**, i.e., computing the exact cost of a solution to a **simplified** version of problem

remove constraints

8-puzzle: Each tile moves independently

simplify problem

8-puzzle: A tile can move to any adjacent position →
Number of moves to get a tile to its goal position = City-Block distance (aka Manhattan distance or L_1 distance)

Map navigation: Use distance formula from n to goal state as $h(n)$, which will always be optimistic because straight line distance is best case, reality has obstacles