

CS-540: Intro to Artificial Intelligence

Informed Search II + Local Search

Lecturer: Erin Winter


A-Search

- Use evaluation function

$$f(n) = g(n) + h(n)$$

where $g(n)$ is minimum cost path from start to current node n (as defined in Dijkstra's)


- Nodes in *the frontier* are sorted by increasing $f(n)$ value
- Think of it as a “first half” $g(n)$ + “second half” $h(n)$ computation



A-Search alone doesn't suffice

- With an arbitrarily bad heuristic function, A-Search isn't **optimal** or **complete**.

A-Search: A,C,G
Cost: 100
Optimal: A,B,C,G
Cost: 3



- We can do better by constraining the heuristic to be **admissible**.
- An admissible heuristic $h(n)$ for all nodes n will never **overestimate** the *true minimum cost path from n to the goal*.

A-Search + Admissible (or
Consistent) Heuristic =

A* Search

effective and used in practice

A* Search Practice

$$f(n) = g(n) + h(n)$$


of nodes tested: 0, expanded: 0

expnd. node	Frontier
	{S:0+8}

To test for admissibility, check that $h(n)$ never overestimates the true cost to the goal.

To check for consistency, check $h(n)$ is consistent AND $h(n) \leq c(n, n') + h(n')$ for all successor nodes n' .

Is h is admissible and/or consistent? It is consistent.




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 1, expanded: 1

expnd. node	Frontier
	{S:8}
S not goal	{A:1+8,B:5+4,C:8+3}




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 2, expanded: 2

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A not goal	{B:9,G:1+9+0,C:11, D:1+3+ ∞ ,E:1+7+ ∞ }




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 3, expanded: 3

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A	{B:9,G:10,C:11,D: ∞ ,E: ∞ }
B not goal	{G:5+4+0, G:10 , C:11, D: ∞ ,E: ∞ } replace




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 4, expanded: 3

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A	{B:9,G:10,C:11,D:∞,E:∞}
B	{G:9,C:11,D:∞,E:∞}
G goal	{C:11,D:∞,E:∞} not expanded




A* Search

$$f(n) = g(n) + h(n)$$

of nodes tested: 4, expanded: 3

expnd. node	Frontier
	{S:8}
S	{A:9,B:9,C:11}
A	{B:9,G:10,C:11,D:∞,E:∞}
B	{G:9,C:11,D:∞,E:∞}
G	{C:11,D:∞,E:∞}



Fast, optimal, complete under same graph conditions as Dijkstra's

path: S,B,G
cost: 9

A* performance in Practice

[from Russell and Norvig, Fig 3.29]

Example
State

1		5
2	6	3
7	4	8

Goal
State

1	2	3
4	5	6
7	8	

For 8-puzzle, average number of states expanded over 100 randomly chosen problems in which optimal path is length ...

... 4 steps

... 8 steps

... 12 steps

Depth-First Iterative Deepening (IDS)

112

6,300

3.6×10^6

A* search using “**number of misplaced tiles**” as the heuristic

13

39

227

A* using “**Sum of City-Block distances**” as the heuristic

12

25

73

Heuristics and A* Performance

- If $h(n) = h^*(n)$ for all n ,
 - only nodes on optimal solution path are expanded
 - no unnecessary work is performed
- If $h(n) = 0$ for all n ,
 - the heuristic is admissible
 - A* performs exactly as Dijkstra's

The closer h is to h^* ,
the fewer extra nodes that will be expanded

How to come up with admissible/consistent heuristics?

Heuristics are often defined by **relaxing the problem**,
i.e., computing the exact cost of a solution to a
simplified version of problem

remove constraints

8-puzzle: Each tile moves independently

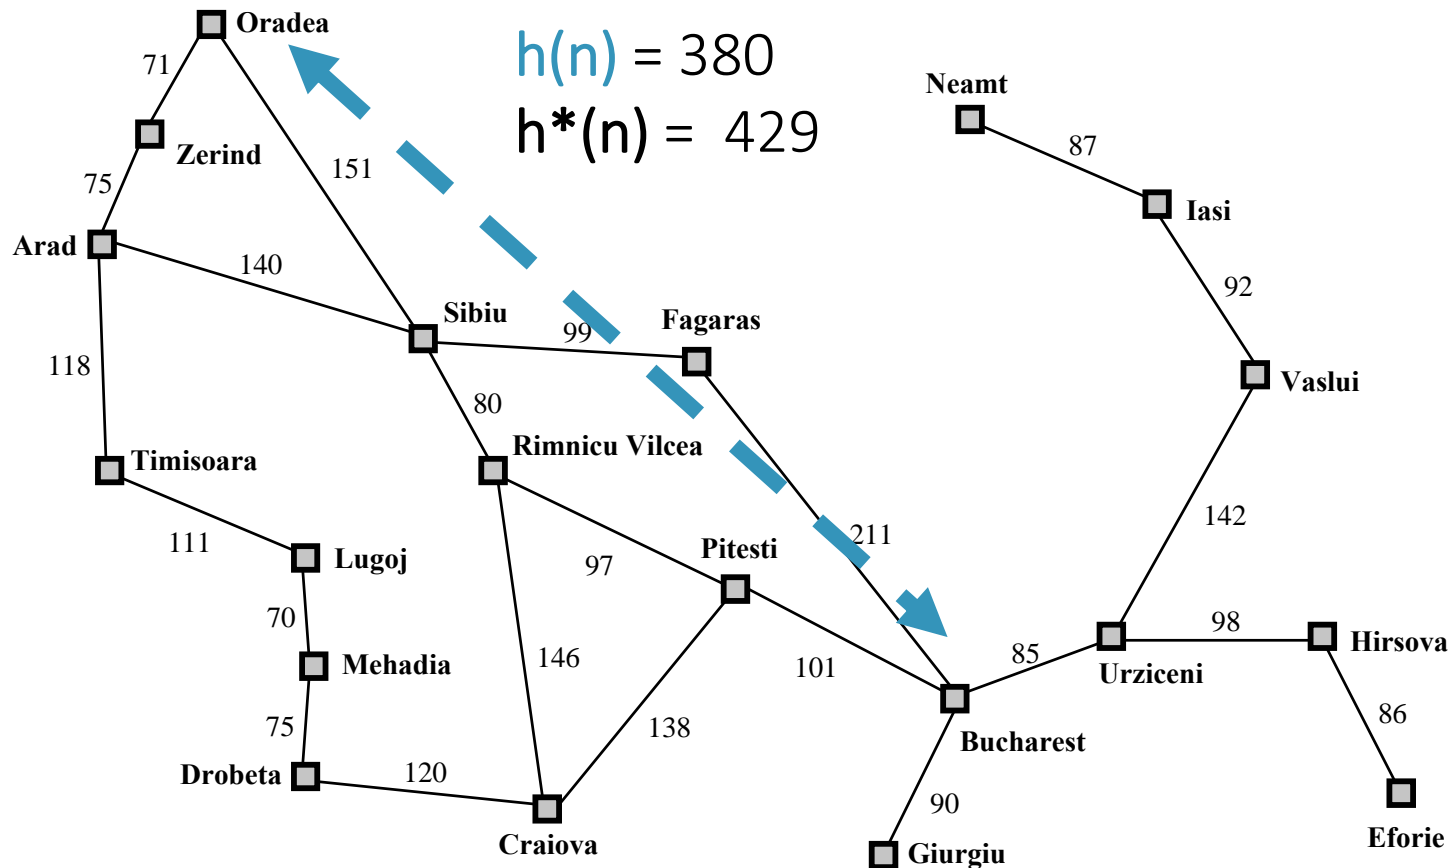
simplify problem

8-puzzle: A tile can move to any adjacent position →

Number of moves to get a tile to its goal position = City-Block distance (aka Manhattan distance or L_1 distance)

Map navigation: Use distance formula from n to goal state as $h(n)$, which will always be optimistic because straight line distance is best case, reality has obstacles

An admissible heuristic is always optimistic. Ex: in this roadmap, $h^*(n)$ from Oradea to Bucharest can't be less than the **straight line distance** from Oradea to Bucharest. Hence, $h(n)$ = straight line distance from Bucharest is admissible.



A* and Heuristics

For an admissible heuristic

- Outside of certain problem domains, h is frequently very simple
- therefore search resorts to (almost) Dijkstra's through parts of the search space
- A* often suffers because it cannot venture down a single path unless it is almost continuously having success (i.e., h is decreasing); any failure to decrease h will cause the search to switch to another path

The Bad News About A^*

- A^* uses a lot of memory, ie. $O(\text{number of states})$
- For some search problems, A^* will fill the available system memory before finding a solution

How do we use less memory?

1. Sacrifice optimality – beam search
2. Sacrifice some time – iterative deepening A^*

Beam Search

Use an evaluation function $f(n) = h(n)$ as in Greedy Best-First search, *and* restrict the maximum size of the Frontier to a constant, k

- Only keep k best nodes as candidates for expansion, and throw away the rest
- More space efficient than Greedy Best-First Search, but may throw away a node on a solution path. Space complexity $O(km)$.
- Not complete
- Not optimal/admissible

Iterating Deepening A*

- Iterative-deepening A*
- Cutoff based on $f(n) = g(n) + h(n)$ value rather than depth
- At each iteration do loop-avoiding DFS, not expanding any node with f -value that exceeds current threshold
- At each iteration increase the f -value threshold by setting it to the smallest f -value of any node that exceeded the cutoff in the previous iteration
- Complete
- Optimal / Admissible
- Linear space required

Informed and Uninformed Search

These methods are for search problems in the following form:

- A directed graph of states S
- Start state(s) $/ \subseteq S$
- Goal state(s) $G \subseteq S$
- A successor function $suc(n)$ that returns all neighbors of state n
- A cost function $c(n, n')$ that returns the cost of moving from state n to state n' *

*This information can be implied through the structure and labels on a graph

Informed and Uninformed Search

- These methods are **systematic**, they search for a **path** from start state to a goal state, then “execute” solution path’s sequence of operators
- Effective on small search spaces
- **not okay for NP-Hard** problems requiring exponential time to find the (optimal) solution

Example NP-Hard Problem: Traveling Salesman

A salesperson wants to visit a list of cities

- stopping in each city only *once*
- (sometimes also must return to the first city)
- goal is to find an ordering of cities that minimizes total distance traveled


Example NP-Hard Problem: Traveling Salesman

A solution is a permutation of cities, called a **tour**

e.g. A – B – C – D – E

assume tours can start at *any* city and do *not* return home

	A	B	C	D	E
A	0	5	8	9	7
B	5	0	6	5	5
C	8	6	0	2	3
D	9	5	2	0	4
E	7	5	3	4	0



Example NP-Hard Problem: Traveling Salesman

Classic NP-Hard problem

How many solutions exist?


$n!$ where n = # of cities

$n = 5$ results in 120 tours

$n = 10$ results in 3,628,800 tours

$n = 20$ results in $\sim 2.4 \times 10^{18}$ tours

	A	B	C	D	E
A	0	5	8	9	7
B	5	0	6	5	5
C	8	6	0	2	3
D	9	5	2	0	4
E	7	5	3	4	0




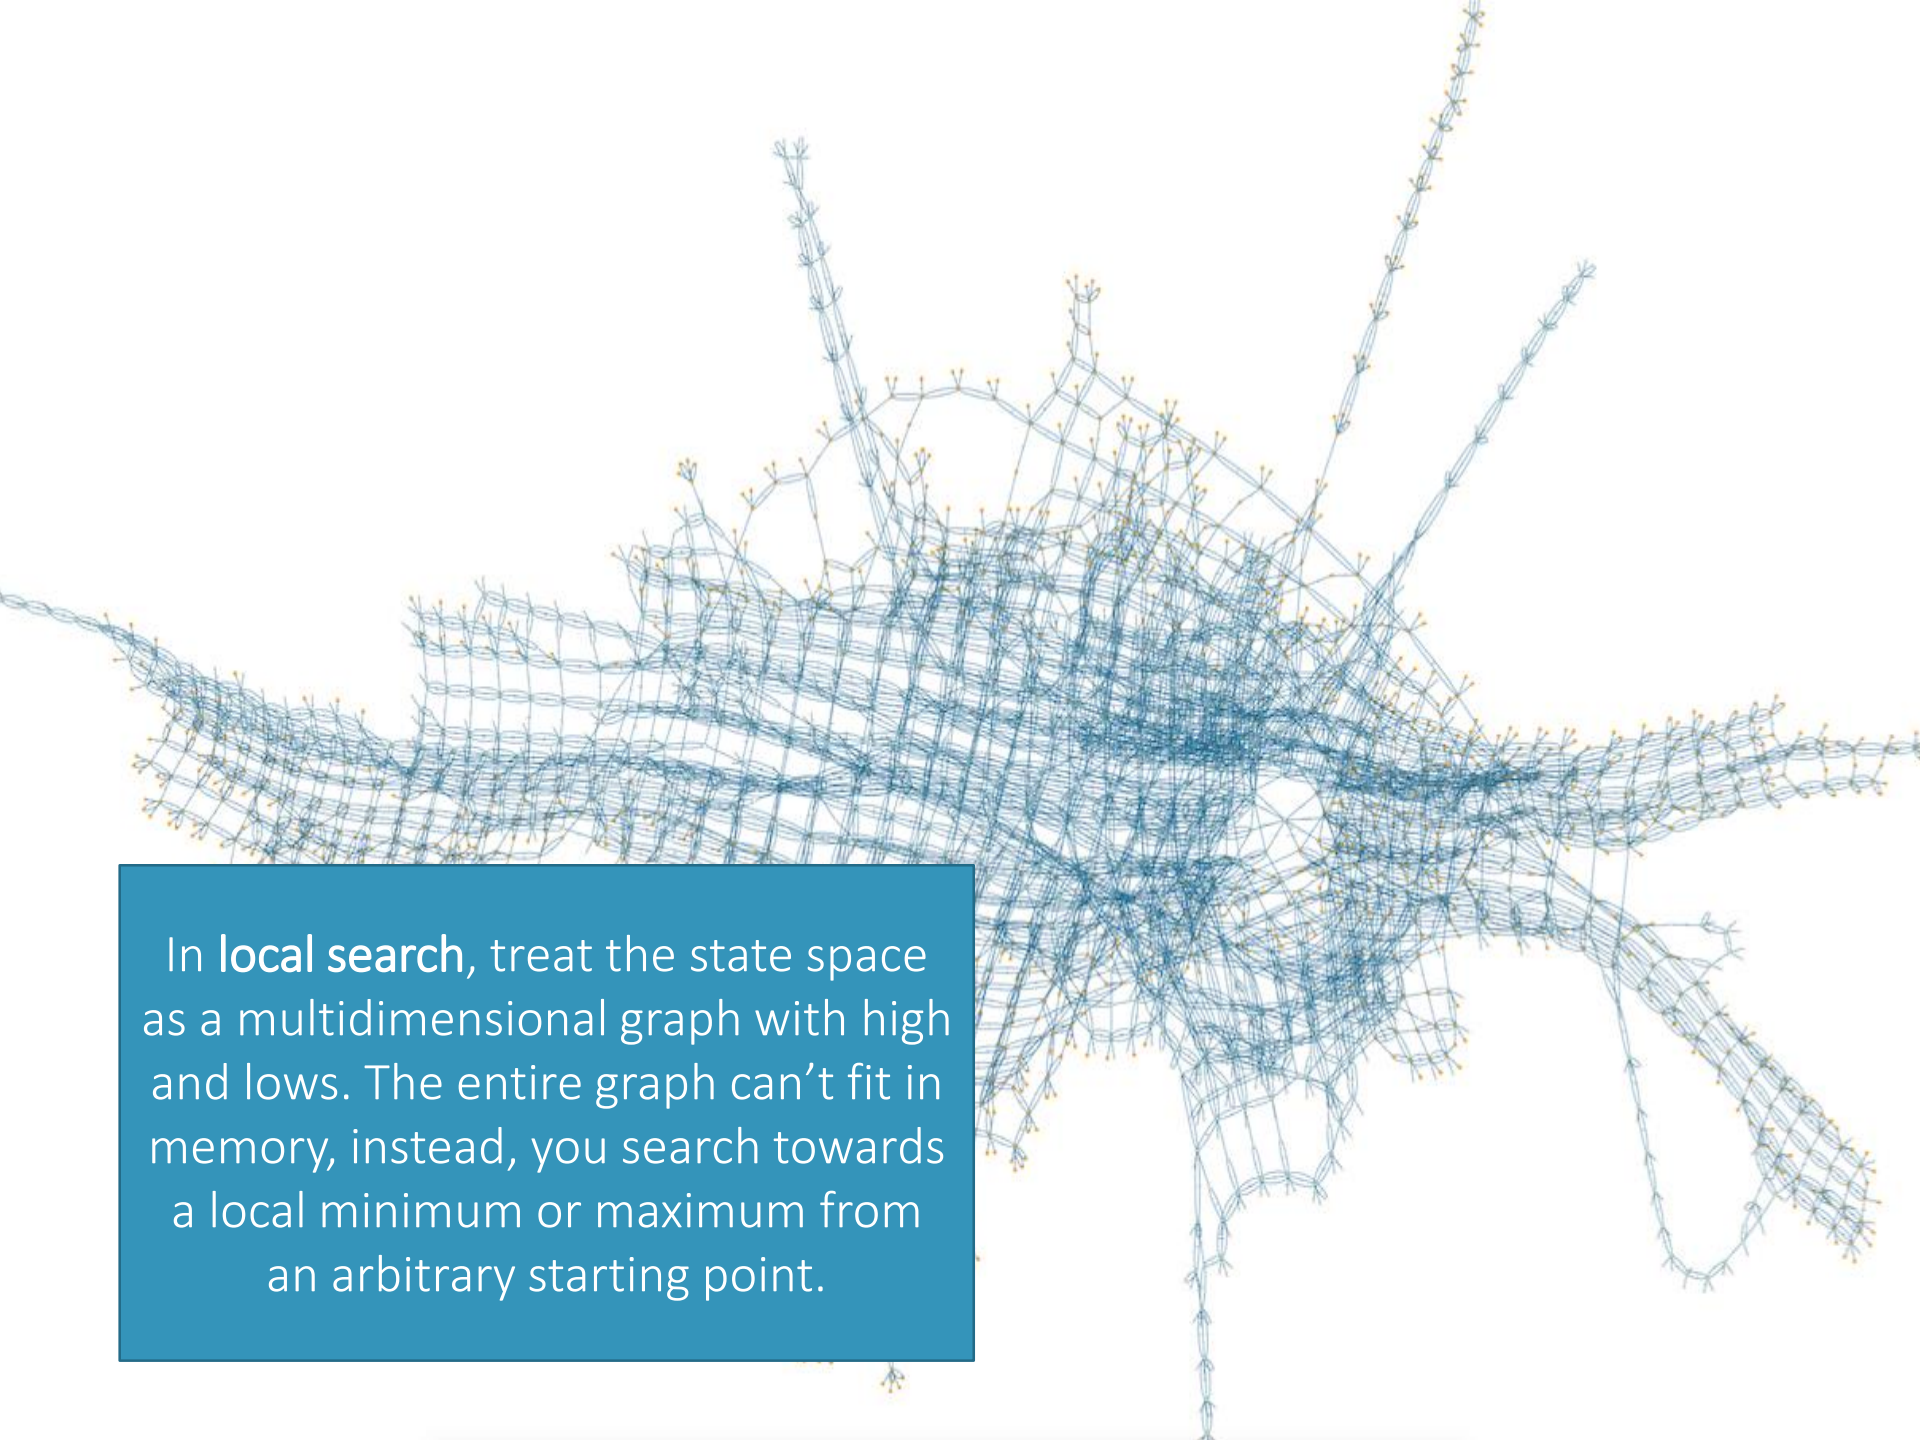
Solving Hard Problems

- Hard problems can be solved in polynomial time by using either an:
 - **approximate model**: find an exact solution to a simpler version of the problem
 - **approximate solution**: find a non-optimal solution to the original hard problem
- We'll explore means to search through a solution space by iteratively improving solutions until one is found that is optimal or near optimal

What if we don't want or need a path? What if just we want a sufficiently **good** state?

Local Search





In **local search**, treat the state space as a multidimensional graph with high and lows. The entire graph can't fit in memory, instead, you search towards a local minimum or maximum from an arbitrary starting point.

Local Search

- Each state n has a **score** $f(n)$ that we can compute
- The goal is to find the **state** (not path) with the highest score, or a reasonably high score
- This is an optimization problem
- Enumerating entire state space is intractable
- Previous search algorithms are too expensive

Local Search

Local searching: every node is a solution

Operators/actions go from one solution to another

Allows you to stop at any time and have a valid solution

Goal: search is to find a *better/best* solution

- No longer searching state space for a solution path and then executing the steps of the solution path
- A* *isn't* a local search since it considers different partial solutions by looking at the estimated cost of a *solution path*

Local Searching

An *operator/action* is needed to transform one solution to another




- TSP: 2-swap operator
 - take two cities and swap their positions in the tour
 - A-B-C-D-E with `swap(A,D)` yields D-B-C-A-E
 - possible since graph is fully connected
- TSP: 2-interchange operator (aka 2-opt swap)
 - reverse the path between two cities
 - A-B-C-D-E with `interchange(A,D)` yields D-C-B-A-E

Neighbors: TSP

state: A-B-C-D-E-F-G-H-A

f = length of tour

2-interchange



Local Searching

Those solutions that can be reached with one application of an operator are in the current solution's *neighborhood* (aka “*move set*”)

Local search considers next only those solutions in the neighborhood

- The neighborhood should be much smaller than the size of the search space (otherwise the search degenerates)

Examples of Neighborhoods

N-queens: Move queen in rightmost, most-conflicting column to a different position in that column

SAT: Flip the assignment of one Boolean variable

Local Searching

An evaluation function, f , is used to map each solution/state to a number corresponding to the *quality/cost* of that solution

- TSP: Use the length of the tour;
A better solution has a shorter tour length
- Maximize f :
called **hill-climbing** (gradient ascent if continuous)
- Minimize f :
called or **valley-finding** (gradient descent if continuous)

Hill-Climbing (HC)

What's a neighbor?

Problem spaces tend to have structure. A *small change* produces a neighboring state

The size of the neighborhood must be small enough for efficiency

Designing the neighborhood is critical; This is the real ingenuity – not the decision to use hill-climbing

Pick which neighbor? The best one (greedy)

What if no neighbor is better than the current state?

Stop

Hill-Climbing Algorithm

1. Pick initial state s
2. Pick t in $\text{neighbors}(s)$ with the largest $f(t)$
3. **if** $f(t) \leq f(s)$ **then** stop and **return** s
4. $s = t$. **Goto** Step 2.

Simple

Greedy

Stops at a *local maximum*

Bonus Section: Proof of A^* Optimality

Proof of A* Optimality (by Contradiction)

- Let

G be the goal in the *optimal* solution

$G2$ be a *sub-optimal* goal found using A* where $f(n) = g(n) + h(n)$, and $h(n)$ is admissible

f^* be the cost of the **optimal path** from Start to G

Hence $g(G2) > f^*$

That is, A* found a sub-optimal path (which it shouldn't)

Proof of A* Optimality (by Contradiction)

- Let n be some node on the *optimal* path but *not* on the path to G_2

- $f(n) \leq f^*$

by admissibility, since $f(n)$ never overestimates the cost to the goal it must be \leq the cost of the optimal path

- $f(G_2) \leq f(n)$

G_2 was chosen over n for the sub-optimal goal to be found

- $f(G_2) \leq f^*$

combining equations

Proof of A* Optimality (by Contradiction)

- $f(G2) \leq f^*$

- $g(G2) + h(G2) \leq f^*$

substituting the definition of f

- $g(G2) \leq f^*$

because $h(G2) = 0$ since $G2$ is a goal node

- This contradicts the assumption that $G2$ was sub-optimal,
 $g(G2) > f^*$

Therefore, A* is optimal with respect to path cost; A* search never finds a sub-optimal goal